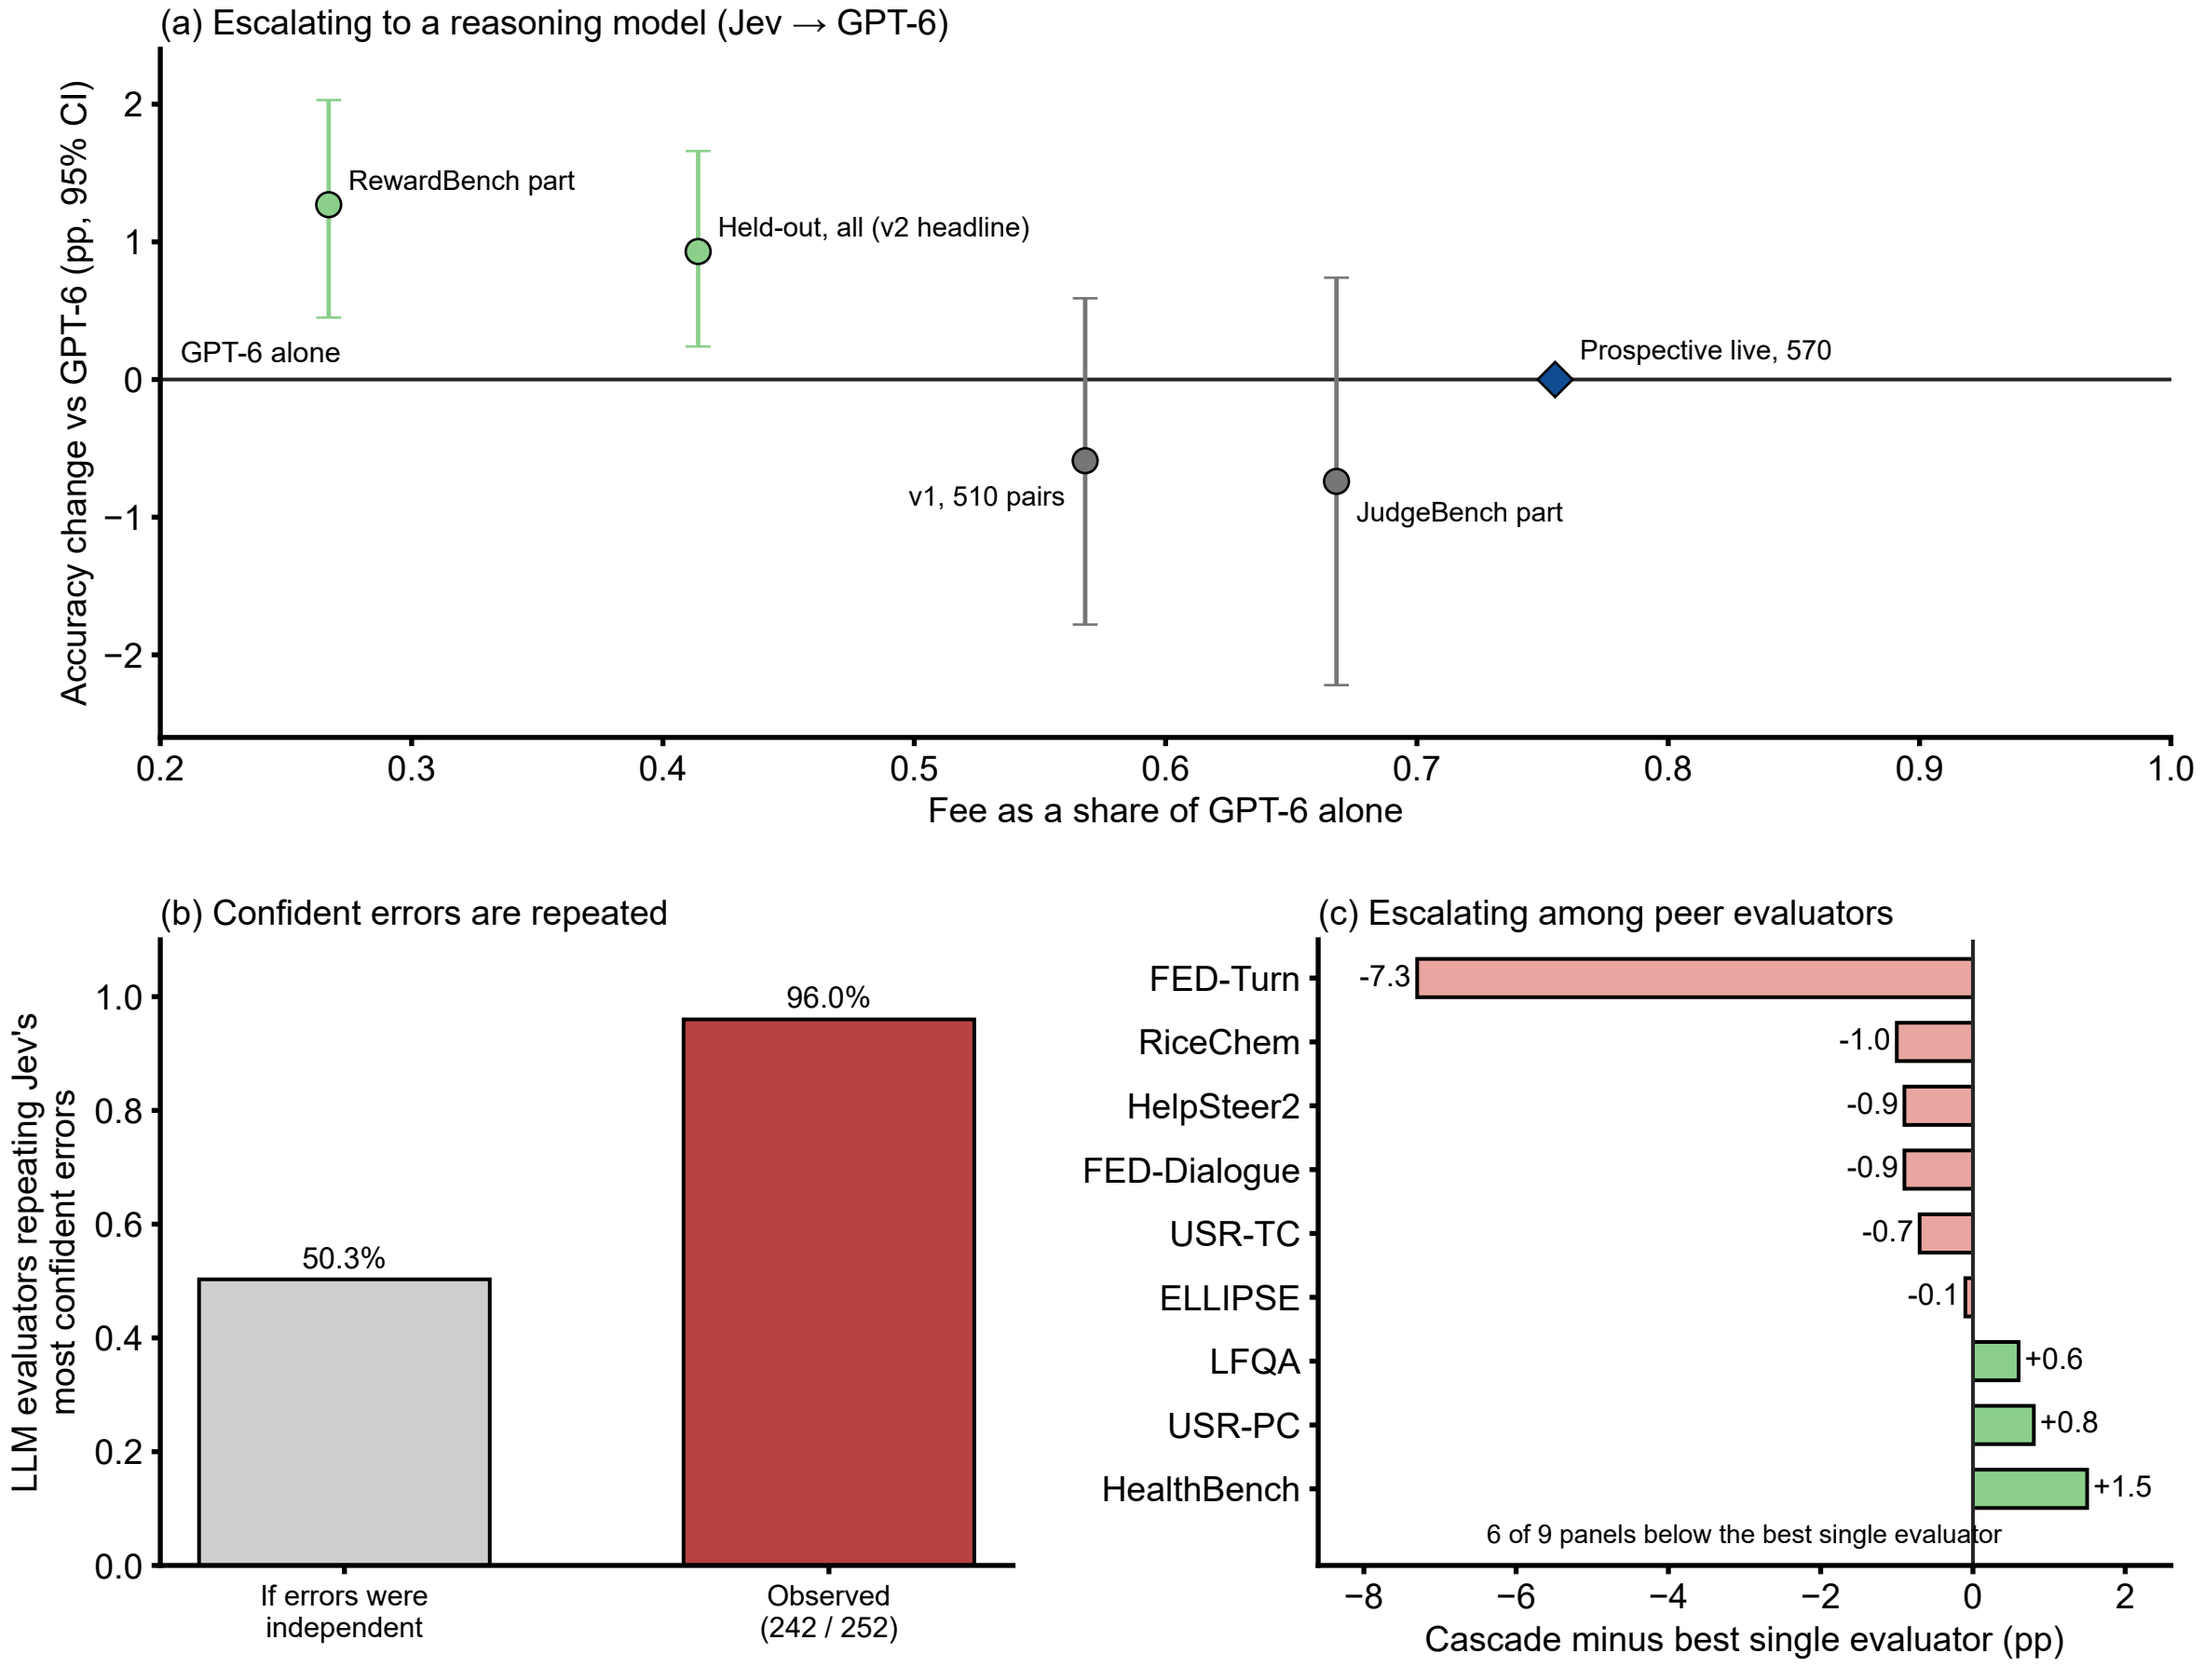
(a) Escalating to a reasoning model (Jev → GPT-6)
Accuracy change vs GPT-6 (pp, 95% CI)
Fee as a share of GPT-6 alone
GPT-6 alone
RewardBench part
Held-out, all (v2 headline)
v1, 510 pairs
JudgeBench part
Prospective live, 570
(b) Confident errors are repeated
LLM evaluators repeating Jev's most confident errors
50.3%
96.0%
If errors were independent
Observed (242 / 252)
(c) Escalating among peer evaluators
FED-Turn
RiceChem
HelpSteer2
FED-Dialogue
USR-TC
ELLIPSE
LFQA
USR-PC
HealthBench
-7.3
-1.0
-0.9
-0.9
-0.7
-0.1
+0.6
+0.8
+1.5
6 of 9 panels below the best single evaluator
Cascade minus best single evaluator (pp)